

# FinalProject\_creditcardfraud

April 19, 2021

## 1 Import necessary libraries for building models

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
import warnings

warnings.filterwarnings("ignore")

dataset = pd.read_csv('creditcard.csv')
RANDOM_STATE = 0
```

## 2 Analyzing the dataset to make decision for building model

```
[2]: # The info show us that there are no null value
# Therefore, we does not have to treat the null value.

dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
 #   Column  Non-Null Count  Dtype  
---  -
 0   Time    284807 non-null float64
 1   V1      284807 non-null float64
 2   V2      284807 non-null float64
 3   V3      284807 non-null float64
 4   V4      284807 non-null float64
 5   V5      284807 non-null float64
 6   V6      284807 non-null float64
 7   V7      284807 non-null float64
 8   V8      284807 non-null float64
 9   V9      284807 non-null float64
```

```

10 V10      284807 non-null float64
11 V11      284807 non-null float64
12 V12      284807 non-null float64
13 V13      284807 non-null float64
14 V14      284807 non-null float64
15 V15      284807 non-null float64
16 V16      284807 non-null float64
17 V17      284807 non-null float64
18 V18      284807 non-null float64
19 V19      284807 non-null float64
20 V20      284807 non-null float64
21 V21      284807 non-null float64
22 V22      284807 non-null float64
23 V23      284807 non-null float64
24 V24      284807 non-null float64
25 V25      284807 non-null float64
26 V26      284807 non-null float64
27 V27      284807 non-null float64
28 V28      284807 non-null float64
29 Amount   284807 non-null float64
30 Class    284807 non-null int64

```

dtypes: float64(30), int64(1)

memory usage: 67.4 MB

```
[3]: dataset.describe()
```

```

[3]:
      count  Time      V1      V2      V3      V4 \
count  284807.000000  2.848070e+05  2.848070e+05  2.848070e+05  2.848070e+05
mean    94813.859575  1.165980e-15  3.416908e-16 -1.373150e-15  2.086869e-15
std     47488.145955  1.958696e+00  1.651309e+00  1.516255e+00  1.415869e+00
min         0.000000 -5.640751e+01 -7.271573e+01 -4.832559e+01 -5.683171e+00
25%     54201.500000 -9.203734e-01 -5.985499e-01 -8.903648e-01 -8.486401e-01
50%     84692.000000  1.810880e-02  6.548556e-02  1.798463e-01 -1.984653e-02
75%    139320.500000  1.315642e+00  8.037239e-01  1.027196e+00  7.433413e-01
max    172792.000000  2.454930e+00  2.205773e+01  9.382558e+00  1.687534e+01

      V5      V6      V7      V8      V9 \
count  2.848070e+05  2.848070e+05  2.848070e+05  2.848070e+05  2.848070e+05
mean    9.604066e-16  1.490107e-15 -5.556467e-16  1.177556e-16 -2.406455e-15
std     1.380247e+00  1.332271e+00  1.237094e+00  1.194353e+00  1.098632e+00
min    -1.137433e+02 -2.616051e+01 -4.355724e+01 -7.321672e+01 -1.343407e+01
25%    -6.915971e-01 -7.682956e-01 -5.540759e-01 -2.086297e-01 -6.430976e-01
50%    -5.433583e-02 -2.741871e-01  4.010308e-02  2.235804e-02 -5.142873e-02
75%     6.119264e-01  3.985649e-01  5.704361e-01  3.273459e-01  5.971390e-01
max     3.480167e+01  7.330163e+01  1.205895e+02  2.000721e+01  1.559499e+01

      ...      V21      V22      V23      V24 \

```

count	...	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
mean	...	1.656562e-16	-3.444850e-16	2.578648e-16	4.471968e-15
std	...	7.345240e-01	7.257016e-01	6.244603e-01	6.056471e-01
min	...	-3.483038e+01	-1.093314e+01	-4.480774e+01	-2.836627e+00
25%	...	-2.283949e-01	-5.423504e-01	-1.618463e-01	-3.545861e-01
50%	...	-2.945017e-02	6.781943e-03	-1.119293e-02	4.097606e-02
75%	...	1.863772e-01	5.285536e-01	1.476421e-01	4.395266e-01
max	...	2.720284e+01	1.050309e+01	2.252841e+01	4.584549e+00

		V25	V26	V27	V28	Amount \
count	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	284807.000000	
mean	5.340915e-16	1.687098e-15	-3.666453e-16	-1.220404e-16	88.349619	
std	5.212781e-01	4.822270e-01	4.036325e-01	3.300833e-01	250.120109	
min	-1.029540e+01	-2.604551e+00	-2.256568e+01	-1.543008e+01	0.000000	
25%	-3.171451e-01	-3.269839e-01	-7.083953e-02	-5.295979e-02	5.600000	
50%	1.659350e-02	-5.213911e-02	1.342146e-03	1.124383e-02	22.000000	
75%	3.507156e-01	2.409522e-01	9.104512e-02	7.827995e-02	77.165000	
max	7.519589e+00	3.517346e+00	3.161220e+01	3.384781e+01	25691.160000	

	Class
count	284807.000000
mean	0.001727
std	0.041527
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	1.000000

[8 rows x 31 columns]

```
[4]: # Check the percentage of fraud and non-fraud transaction
non_fraud_trans = dataset['Class'].value_counts()[0] # number of non-fraud
↳ transaction
fraud_trans = dataset['Class'].value_counts()[1] # number of fraud transaction
total_trans = len(dataset) # total number of transactions in the whole datasets
print('No Frauds =', non_fraud_trans, "which equal", round(non_fraud_trans/
↳ total_trans * 100, 2), '% of the dataset')
print('Frauds =', fraud_trans, "which equal", round(fraud_trans/total_trans *
↳ 100,2), '% of the dataset')


# The result show us that the dataset for credit card fraud is really
↳ imbalanced
# because we have 99.83% of the dataset is non-fraudulent
# and only 0.17% of the dataset is fradulent transaction
```

No Frauds = 284315 which equal 99.83 % of the dataset

Frauds = 492 which equal 0.17 % of the dataset


```
[5]: plt.figure(figsize=(18,4))
plt.title('Distribution of Time', fontsize = 16)
sns.distplot(dataset.Time, color = 'r')
```

```
[5]: <AxesSubplot:title={'center':'Distribution of Time'}, xlabel='Time',
ylabel='Density'>
```



```
[6]: plt.figure(figsize=(18,4))
plt.title('Distribution of Amount')
sns.distplot(dataset.Amount, color = 'r')
```

```
[6]: <AxesSubplot:title={'center':'Distribution of Amount'}, xlabel='Amount',
ylabel='Density'>
```



### 3 Preprocessing Technique

1. Scaling and Distributing
2. Splitting dataset

```
[7]: # Scaling and Distributing
# Since the whole dataset is already scaled except for the time and amount
```

```
# thus we have to scale the time and amount columns of the dataset before we
↳working on building models
```

```
std_scaler = StandardScaler()
```

```
dataset.insert(0, 'scaled_amount', std_scaler.fit_transform(dataset['Amount'].
↳values.reshape(-1,1)))
```

```
dataset.insert(1, 'scaled_time', std_scaler.fit_transform(dataset['Time'].
↳values.reshape(-1,1)))
```

```
[8]: dataset.drop(['Amount','Time'], axis = 1, inplace = True)
```

```
[9]: dataset.head() # Use this to make visualize the dataset after the modification
↳above
```

```
[9]:
```

	scaled_amount	scaled_time	V1	V2	V3	V4	\
0	0.244964	-1.996583	-1.359807	-0.072781	2.536347	1.378155	
1	-0.342475	-1.996583	1.191857	0.266151	0.166480	0.448154	
2	1.160686	-1.996562	-1.358354	-1.340163	1.773209	0.379780	
3	0.140534	-1.996562	-0.966272	-0.185226	1.792993	-0.863291	
4	-0.073403	-1.996541	-1.158233	0.877737	1.548718	0.403034	

	V5	V6	V7	V8	...	V20	V21	V22	\
0	-0.338321	0.462388	0.239599	0.098698	...	0.251412	-0.018307	0.277838	
1	0.060018	-0.082361	-0.078803	0.085102	...	-0.069083	-0.225775	-0.638672	
2	-0.503198	1.800499	0.791461	0.247676	...	0.524980	0.247998	0.771679	
3	-0.010309	1.247203	0.237609	0.377436	...	-0.208038	-0.108300	0.005274	
4	-0.407193	0.095921	0.592941	-0.270533	...	0.408542	-0.009431	0.798278	

	V23	V24	V25	V26	V27	V28	Class
0	-0.110474	0.066928	0.128539	-0.189115	0.133558	-0.021053	0
1	0.101288	-0.339846	0.167170	0.125895	-0.008983	0.014724	0
2	0.909412	-0.689281	-0.327642	-0.139097	-0.055353	-0.059752	0
3	-0.190321	-1.175575	0.647376	-0.221929	0.062723	0.061458	0
4	-0.137458	0.141267	-0.206010	0.502292	0.219422	0.215153	0

```
[5 rows x 31 columns]
```

```
[10]: X = dataset.iloc[:, dataset.columns != 'Class']
y = dataset['Class']
```

```
[11]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.4,
↳random_state = RANDOM_STATE)

X_test, X_val, y_test, y_val = train_test_split(X_test, y_test, test_size = 0.
↳5, random_state = RANDOM_STATE)
```

#### 4 Use simple parameter for Learning Algorithm on the imbalance dataset to prove that the model will be overfitting if we don't use the resample training set

```
[12]: from sklearn.linear_model import LogisticRegression
      from sklearn.ensemble import RandomForestClassifier

      # Use Logistic Regression on original dataset
      lr = LogisticRegression(random_state = RANDOM_STATE).fit(X_train, y_train)
      y_lr_pred = lr.predict(X_test)

      # User Random Forest on undersample dataset
      rfc = RandomForestClassifier(random_state = RANDOM_STATE).fit(X_train, y_train)
      y_rf_pred = rfc.predict(X_test)
```

#### 5 Use confusion matrix to visualize the original dataset

```
[13]: from sklearn.metrics import confusion_matrix

      # Logistic Regression confusion matrix
      lr_cm = confusion_matrix(y_test, y_lr_pred)

      # Random Forest confusino matrix
      rf_cm = confusion_matrix(y_test, y_rf_pred)
```


```
[14]: fig= plt.figure(figsize=(15,10))

      ax1 = fig.add_subplot(2,2,1)
      sns.heatmap(lr_cm, cmap="coolwarm_r",annot=True, fmt = "", linewidths=0.5)
      plt.title("Logistic Regression")
      plt.xlabel("Predicted class")
      plt.ylabel("Real class")

      ax2 = fig.add_subplot(2,2,2)
      sns.heatmap(rf_cm, cmap="coolwarm_r",annot=True, fmt = "", linewidths=0.5)
      plt.title("Random Forest Classifer")
      plt.xlabel("Predicted class")
      plt.ylabel("Real class")

      plt.show()

      # TN = 0,0 no. of normal transaction which are predited normal
      # FP = 0,1 no of normal transaction which are predicted fraud
      # FN = 1,0 no of fraud Transaction which are predicted normal
      # TP = 1,1 no of fraud transaction which are predicted fraud
```



```
[15]: from sklearn.metrics import classification_report

print("\t\t Overfitting Log Regression \n\n",classification_report(y_test,
    ↪y_lr_pred))

print("\t\t Overfitting RF \n\n",classification_report(y_test, y_rf_pred))
```

#### Overfitting Log Regression

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56871
1	0.91	0.58	0.71	90
accuracy			1.00	56961
macro avg	0.96	0.79	0.85	56961
weighted avg	1.00	1.00	1.00	56961

#### Overfitting RF

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56871
1	0.97	0.76	0.85	90
accuracy			1.00	56961
macro avg	0.99	0.88	0.92	56961
weighted avg	1.00	1.00	1.00	56961

```
[16]: # This is still preprocessing dataset
from imblearn.under_sampling import RandomUnderSampler
```

```
undersample = RandomUnderSampler(sampling_strategy='majority')
X_train_us, y_train_us= undersample.fit_resample(X_train, y_train)
```

## 6 Use GridSearchCV to find the best parameters for each classifier on the resample dataset

1. Find the best parameters for Logistic Regression on resampled dataset

```
[17]: from sklearn.model_selection import GridSearchCV

lr_params = {"penalty": ['l1', 'l2'],
             'C': [0.001, 0.01, 0.1, 1, 10, 100, 1000]}

grid_lr = GridSearchCV(LogisticRegression(), lr_params)

grid_lr.fit(X_train_us, y_train_us)

[17]: GridSearchCV(estimator=LogisticRegression(),
                  param_grid={'C': [0.001, 0.01, 0.1, 1, 10, 100, 1000],
                              'penalty': ['l1', 'l2']})
```

```
[18]: grid_lr.best_params_
```

```
[18]: {'C': 0.01, 'penalty': 'l2'}
```

2. Find the best parameters for Support Vector Machine on original dataset and resampled dataset

```
[19]: rf_params = {'n_estimators': [8, 16, 64, 200],
                  'max_depth': [None, 10, 50, 200, 1000]}

grid_rf = GridSearchCV(RandomForestClassifier(), rf_params)

grid_rf.fit(X_train_us, y_train_us)

[19]: GridSearchCV(estimator=RandomForestClassifier(),
                  param_grid={'max_depth': [None, 10, 50, 200, 1000],
                              'n_estimators': [8, 16, 64, 200]})
```

```
[20]: grid_rf.best_params_
```

```
[20]: {'max_depth': 1000, 'n_estimators': 200}
```

## 7 Fit the dataset with the classifier

1. Fit train dataset on Logistic Regression Classifier



```
[21]: lr = LogisticRegression(C = 0.01, penalty = 'l2', random_state=RANDOM_STATE)
lr.fit(X_train_us, y_train_us)
y_lr_pred = lr.predict(X_test)
```

2. Fit train dataset on Random Forest

```
[22]: rf = RandomForestClassifier(max_depth = None, n_estimators = 16)
rf.fit(X_train_us, y_train_us)
y_rf_pred = rf.predict(X_test)
```

```
[40]: from sklearn.model_selection import cross_val_score

lr_cv_score = cross_val_score(lr, X_train_us, y_train_us, cv=5)

print("Logistic Regression has",round(lr_cv_score.mean(), 2) * 100, "% cross_
    ↳validation score")

rf_cv_score = cross_val_score(rf, X_train_us, y_train_us, cv=5)

print("Random Forest Classifier has",round(rf_cv_score.mean(), 2) * 100, "%_
    ↳cross validation score")
```

Logistic Regression has 94.0 % cross validation score

Random Forest Classifier has 93.0 % cross validation score

## 8 Use confusion matrix for visualization

```
[24]: from sklearn.metrics import confusion_matrix

# Logistic Regression confusion matrix
lr_cm = confusion_matrix(y_test, y_lr_pred)

# Support Vector Classifier confusion matrix
rf_cm = confusion_matrix(y_test, y_rf_pred)
```

```
[44]: lr_cm
```

```
[44]: array([[56213,   658],
           [   14,    76]], dtype=int64)
```

```
[45]: rf_cm
```

```
[45]: array([[55287,  1584],
           [   13,    77]], dtype=int64)
```

```
[42]: print("\t\t Resampled Log Regression \n\n",classification_report(y_test,_
    ↳y_lr_pred))
```

```
print("\t\t Resampled Random Forest \n\n",classification_report(y_test,
→y_rf_pred))
```

#### Resampled Log Regression

	precision	recall	f1-score	support
0	1.00	0.99	0.99	56871
1	0.10	0.84	0.18	90
accuracy			0.99	56961
macro avg	0.55	0.92	0.59	56961
weighted avg	1.00	0.99	0.99	56961

#### Resampled Random Forest

	precision	recall	f1-score	support
0	1.00	0.97	0.99	56871
1	0.05	0.86	0.09	90
accuracy			0.97	56961
macro avg	0.52	0.91	0.54	56961
weighted avg	1.00	0.97	0.98	56961

```
[51]: # Let's Plot LogisticRegression Learning Curve
from sklearn.model_selection import ShuffleSplit
from sklearn.model_selection import learning_curve

def plot_learning_curve(estimator1, estimator2, X, y, ylim=None, cv=None,
→n_jobs=1, train_sizes=np.linspace(.1, 1.0, 5)):

    f, (ax1, ax2) = plt.subplots(1,2, figsize=(20,10), sharey=True)

    if ylim is not None:
        plt.ylim(*ylim)

    # First Estimator
    train_sizes, train_scores, test_scores = learning_curve(estimator1, X, y,
→cv=cv, n_jobs=n_jobs, train_sizes=train_sizes)

    ax1.plot(train_sizes, np.mean(train_scores, axis=1), 'o-', color="#ff9124",
→label="Training score")
    ax1.plot(train_sizes, np.mean(test_scores, axis=1), 'o-', color="#2492ff",
→label="Cross-validation score")
```

```

ax1.set_title("Logistic Regression Learning Curve", fontsize=14)
ax1.set_xlabel('Training size (m)')
ax1.set_ylabel('Score')
ax1.grid(True)
ax1.legend(loc="best")

# Second Estimator
train_sizes, train_scores, test_scores = learning_curve(estimator2, X, y,
cv=cv, n_jobs=n_jobs, train_sizes=train_sizes)

ax2.plot(train_sizes, np.mean(train_scores, axis=1), 'o-', color="#ff9124",
label="Training score")
ax2.plot(train_sizes, np.mean(test_scores, axis=1), 'o-', color="#2492ff",
label="Cross-validation score")

ax2.set_title("Random Forest Learning Curve", fontsize=14)
ax2.set_xlabel('Training size (m)')
ax2.set_ylabel('Score')
ax2.grid(True)
ax2.legend(loc="best")

return plt

```

```

[52]: cv = ShuffleSplit(n_splits=100, test_size=0.2, random_state=0)



plot_learning_curve(lr, rf, X_train_us, y_train_us, (0.87, 1.01), cv=cv,
n_jobs=2)

```

```

[52]: <module 'matplotlib.pyplot' from 'D:\\Python\\lib\\site-
packages\\matplotlib\\pyplot.py'>

```



```
[ ]:
```

```
[29]: import tensorflow as tf
from tensorflow import keras
from keras import backend as K
from keras.models import Sequential
from keras.layers import Activation
from keras.layers.core import Dense
from keras.optimizers import Adam
from keras.metrics import categorical_crossentropy

n_inputs = X_train_us.shape[1]

nn_model = keras.models.Sequential([
    keras.layers.Dense(n_inputs, input_shape=(n_inputs, ), activation='relu'),
    keras.layers.Dense(32, activation = "relu"),
    keras.layers.Dense(2, activation = "softmax")
])
```

```
[30]: nn_model.compile(optimizer="adam", loss='sparse_categorical_crossentropy',
    ↪metrics=['accuracy'])
```

```
[32]: epoch = 20

history = nn_model.fit(X_train_us, y_train_us, validation_data=(X_val,y_val),
    ↪batch_size=25, epochs = epoch, shuffle=True, verbose=2)
```

Epoch 1/20

WARNING:tensorflow:Callbacks method `on\_test\_batch\_end` is slow compared to the batch time (batch time: 0.0000s vs `on\_test\_batch\_end` time: 0.0010s). Check your callbacks.

24/24 - 2s - loss: 0.5391 - accuracy: 0.5956 - val\_loss: 0.6786 - val\_accuracy: 0.6388

Epoch 2/20

24/24 - 2s - loss: 0.3445 - accuracy: 0.8771 - val\_loss: 0.4628 - val\_accuracy: 0.9192

Epoch 3/20

24/24 - 2s - loss: 0.2622 - accuracy: 0.9437 - val\_loss: 0.3296 - val\_accuracy: 0.9759

Epoch 4/20

24/24 - 2s - loss: 0.2143 - accuracy: 0.9488 - val\_loss: 0.2621 - val\_accuracy: 0.9807

Epoch 5/20


24/24 - 2s - loss: 0.1786 - accuracy: 0.9522 - val\_loss: 0.2003 - val\_accuracy: 0.9835

Epoch 6/20  
24/24 - 2s - loss: 0.1560 - accuracy: 0.9522 - val\_loss: 0.1696 - val\_accuracy: 0.9815  
Epoch 7/20  
24/24 - 2s - loss: 0.1385 - accuracy: 0.9522 - val\_loss: 0.1516 - val\_accuracy: 0.9794  
Epoch 8/20  
24/24 - 2s - loss: 0.1249 - accuracy: 0.9539 - val\_loss: 0.1417 - val\_accuracy: 0.9777  
Epoch 9/20  
24/24 - 2s - loss: 0.1128 - accuracy: 0.9573 - val\_loss: 0.1401 - val\_accuracy: 0.9735  
Epoch 10/20  
24/24 - 2s - loss: 0.1041 - accuracy: 0.9608 - val\_loss: 0.1278 - val\_accuracy: 0.9744  
Epoch 11/20  
24/24 - 2s - loss: 0.0965 - accuracy: 0.9676 - val\_loss: 0.1205 - val\_accuracy: 0.9733  
Epoch 12/20  
24/24 - 2s - loss: 0.0901 - accuracy: 0.9693 - val\_loss: 0.1163 - val\_accuracy: 0.9719  
Epoch 13/20  
24/24 - 2s - loss: 0.0845 - accuracy: 0.9710 - val\_loss: 0.1231 - val\_accuracy: 0.9686  
Epoch 14/20  
24/24 - 2s - loss: 0.0781 - accuracy: 0.9727 - val\_loss: 0.1064 - val\_accuracy: 0.9731  
Epoch 15/20  
24/24 - 2s - loss: 0.0742 - accuracy: 0.9727 - val\_loss: 0.1186 - val\_accuracy: 0.9674  
Epoch 16/20  
24/24 - 2s - loss: 0.0700 - accuracy: 0.9761 - val\_loss: 0.1167 - val\_accuracy: 0.9671  
Epoch 17/20  
24/24 - 2s - loss: 0.0659 - accuracy: 0.9795 - val\_loss: 0.1175 - val\_accuracy: 0.9661  
Epoch 18/20  
24/24 - 2s - loss: 0.0620 - accuracy: 0.9795 - val\_loss: 0.1174 - val\_accuracy: 0.9660  
Epoch 19/20  
24/24 - 1s - loss: 0.0582 - accuracy: 0.9846 - val\_loss: 0.1090 - val\_accuracy: 0.9687  
Epoch 20/20  
24/24 - 1s - loss: 0.0553 - accuracy: 0.9846 - val\_loss: 0.1196 - val\_accuracy: 0.9649

```
[33]: fig, (ax1, ax2) = plt.subplots(2, 1, figsize=(12, 12))
ax1.plot(history.history['loss'], color='b', label="Training loss")
ax1.plot(history.history['val_loss'], color='r', label="validation loss")
ax1.set_xticks(np.arange(1, epoch, 1))
ax1.set_yticks(np.arange(0, 1, 0.1))
ax1.legend(loc='best', shadow=True)


ax2.plot(history.history['accuracy'], color='b', label="Training accuracy")
ax2.plot(history.history['val_accuracy'], color='r', label="Validation accuracy")
ax2.set_xticks(np.arange(1, epoch, 1))
ax2.legend(loc='best', shadow=True)

plt.tight_layout()
plt.show()
```



```
[34]: pd.DataFrame(history.history).plot(figsize=(10,5))
plt.grid(True)
plt.gca().set_ylim(0,1)
```

[34]: (0.0, 1.0)



```
[41]: history.model.evaluate(X_test, y_test)
```

1781/1781 [=====] - 1s 841us/step - loss: 0.1160 -  
accuracy: 0.9660

[41]: [0.11600951850414276, 0.9659942984580994]